



# PostGIS & MongoDB

Feature & Performance Comparison for Geospatial Data Model


# Introduction & Data

## Hypothesis:

Since PostGIS is built specifically for geospatial data, it exceeds MongoDB in performance when using geospatial data.

## Primary Datasets used for Testing:

1. Sydney Airbnb Open Data
2. OpenStreetMap Australia




# Agenda

- ❏ Spherical Geospatial Functions
- ❏ Intersects & Within Geospatial Functions
- ❏ Dimensionally Extended 9 Intersection Model in PostGIS

## # Feature 1: \$nearSphere in MongoDB vs ST\_DWithin in PostGIS

### Question:

How many Airbnb listings are near Bondi Beach by setting up with different maximum distance and data size?



Bondi Beach Coordinates:  
[151.2767, -33.8915]



### Maximum Distance:

- ❑ 5 kilometers
- ❑ 10 kilometers
- ❑ 20 kilometers
- ❑ 30 kilometers

### Data Points from Airbnb listings:


- ❑ 100 Data Points
- ❑ 1000 Data Points
- ❑ 10,000 Data Points
- ❑ 100,000 Data Potins
- ❑ 1 million Data Points

\*each query will be run by ten times

## # Feature 1: Time Execution Statistics Table

	MongoDB: \$nearSphere		PostGIS : ST_Dwithin	
Data Size/max distance (km)	mean time [second]	standard deviation	mean time [second]	standard deviation
100				
5	0.007187605	0.002276653	0.043274999	0.006481027
10	0.008389878	0.005950478	0.043648839	0.00233995
20	0.010199976	0.002533084	0.048395753	0.007962633
30	0.008744836	0.002221548	0.045145726	0.002509106
1k				
5	0.025807548	0.016310987	0.065703702	0.022447455
10	0.035507584	0.010341971	0.081143594	0.027774079
20	0.040365314	0.017878762	0.082090163	0.015813456
30	0.03627193	0.004936658	0.103290701	0.031626384
10k				
5	0.141123176	0.075665804	1.241498327	0.325197246
10	0.241815591	0.029083681	2.782282567	0.201445733
20	0.301468277	0.03964041	3.650822353	0.30310949
30	0.475698066	0.444679669	4.069888043	0.407288217
100k				
5	1.089699531	0.856567986	1.241498327	0.325197246
10	2.787020278	0.520455389	2.782282567	0.201445733
20	4.641707802	0.674566761	3.650822353	0.30310949
30	6.599097824	1.361575328	4.069888043	0.407288217
1 million				
5	34.98340652	31.24455561	74.07140317	25.09023428
10	317.8399241	272.6782824	75.91673875	8.069031153
20	914.9281501	353.0617526	107.6204711	7.79144199
30	1149.652263	677.2832733	341.3338144	719.1305923

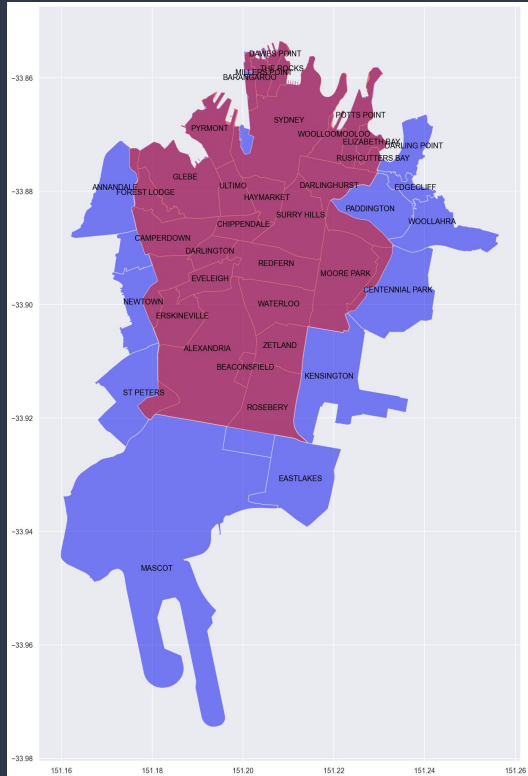
## # Feature 1: Comparison of Time Performance






## # Feature 2

Polygon Filter (red)

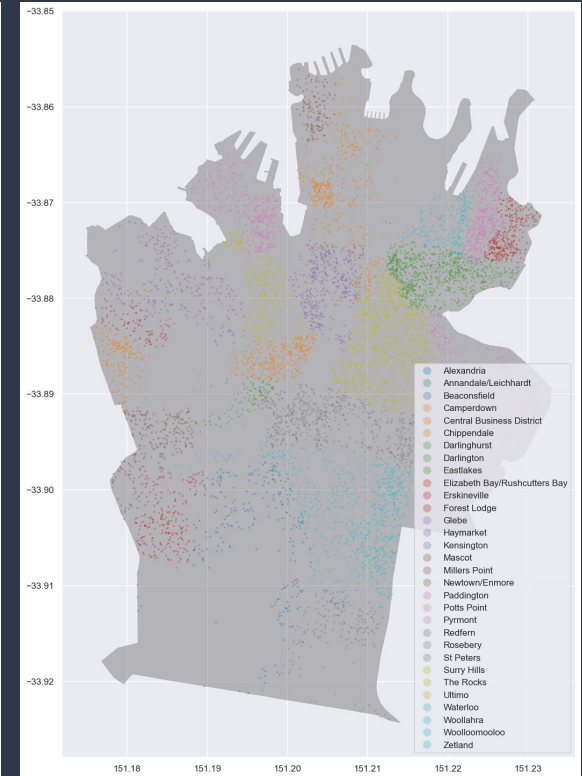


>>>

Intersects (with roads)




Within (Airbnb listings)




# Performance Comparison

Intersects() Runtime Comparison between PostGIS and MongoDB



Within() Runtime Comparison between PostGIS and MongoDB





# Execution stats

## Intersects

system	time in milliseconds			
	PostGIS		MongoDB	
	mean	std	mean	std
n				
100	24.532204	3.141809	19.28	7.039365
1000	24.205370	2.090392	30.84	7.220605
10000	28.385696	3.764740	106.34	34.110630
100000	101.556478	20.511965	1061.42	106.424274
500000	291.452546	69.281944	4414.42	1021.952714

## Within

system	time in milliseconds			
	PostGIS		MongoDB	
	mean	std	mean	std
n				
100	25.001583	2.027436	23.72	4.965349
1000	30.667968	2.160910	39.16	5.056255
10000	87.840109	19.978261	149.86	45.633862
100000	772.345309	334.494025	1130.92	441.380711
1000000	12163.742747	8471.295160	11361.42	2825.621074



## # Feature 3: Dimensionally Extended 9 Intersection Model

- Finding specific types of intersecting roads:
  - Cross intersections: 0F1FF0F02



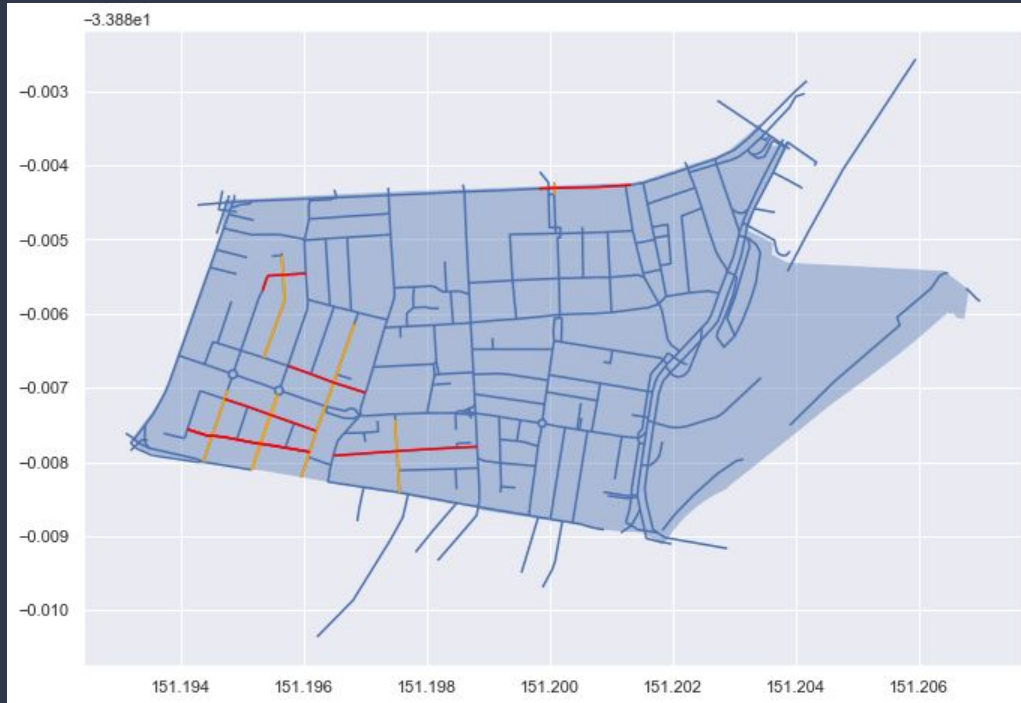
		Road Y X-----X		
		Interior	Boundary	Exterior
Road X X     X	Interior	0	F	1
	Boundary	F	F	0
	Exterior	1	0	2

- T intersections: F010F0102



		Road Y X-----X		
		Interior	Boundary	Exterior
Road X X     X	Interior	F	0	1
	Boundary	0	F	0
	Exterior	1	0	2

## # Feature 3: Dimensionally Extended 9 Intersection Model



Cross Intersections in Chippendale

- Small number of cross intersections found and highlighted

## # Feature 3: Dimensionally Extended 9 Intersection Model



Cross Intersections in Alexandria

- None found in Chippendale
- 2 were found in Alexandria

## # Feature 3: Dimensionally Extended 9 Intersection Model

- Roads composed of multiple segments/records.
- Feature works as intended.
- Effectiveness determined by data model structure.
- Recommend pre-processing dataset for increasing number of results.

	osm_id	code	fclass	name	ref	oneway	maxspeed	layer	bridge	tunnel	geometry
7	411192715	5113	primary	Cleveland Street	None	B	50	0	F	F	LINESTRING (151.19981 -33.88874, 151.20001 -33...
12	411017206	5113	primary	Cleveland Street	None	F	50	0	F	F	LINESTRING (151.19349 -33.88781, 151.19426 -33...
19	411192716	5113	primary	Cleveland Street	None	B	50	0	F	F	LINESTRING (151.20037 -33.88882, 151.20049 -33...
23	595889192	5113	primary	Cleveland Street	None	B	50	0	F	F	LINESTRING (151.19822 -33.88851, 151.19849 -33...
40	175969863	5113	primary	Cleveland Street	None	B	50	0	F	F	LINESTRING (151.19516 -33.88811, 151.19507 -33...



# Conclusion

## MongoDB:

- ❖ Ease of setup (or the lack thereof)
- ❖ Tends to outperform PostGIS when using spherical filters (near) if maximum radius is less than 100,000km
- ❖ May outperform PostGIS when using spatial filters (Within) if two conditions are met
  - Only dealing with points
  - Small dataset (< 1000 data points)
- ❖ Equipped with only essential functions
  - Near, Within, Intersects
- ❖ Does not support
  - Coordinate reprojection (projected to non-projected, vice versa)
  - Geometry transformation (to\_centroid, buffer, unary\_union)
  - Distance calculation
- ❖ Due to limitation of geospatial features, often requires use of programming language such as Python to perform necessary transformations
  - Converting all geometries to the same coordinate system
  - Geometric manipulations

## PostGIS:

- ❖ Ease of setup
- ❖ Tends to outperform MongoDB when
  - Dealing with multiple geometries, using lines and polygons
  - Dataset is large (> 1000 data points)
- ❖ Better performance consistency - standard deviation of execution time is usually very small
- ❖ Comprehensive and powerful support for geospatial data
  - Coordinate reprojection
  - Geometric manipulation / transformation
    - Buffer, centroid, distance calculation, unary union, aggregation and merge of geometry
  - Syntactic difference makes it easy to chain geospatial functions compared to MongoDB
  - DE-9IM for specific intersection projection or filtering
- ❖ Usually comes with GDAL and OGR when installing PostGIS
  - GDAL for raster data import / export / conversion
  - OGR for vector data import / export / conversion

**Thank you!**

# Sample Queries: \$nearSphere vs ST\_DWithin

```
from bson.son import SON # required for $maxDistance
def distance_to_radians(km):
    return km/6378.137

#bondi beach coordinates
lat = -33.8915
lon = 151.2767
point= [lon, lat]

km=5
result= db.airbnb_sydney_100.find(
    {
        "geometry": SON([
            ("nearSphere", point),
            ("maxDistance", distance_to_radians(km))
        ])
    }
)
results= [item for item in result]
print(len(results))
```

**MongoDB : \$nearSphere + \$maxdistance**

```
sql=''
SELECT *
FROM airbnb_sydney_100
WHERE ST_DWithin(airbnb_sydney_100.geometry,
ST_SetSRID(ST_MakePoint(151.2767,-33.8915),4326),0.05)'''

result =gpd.read_postgis(sql, engine, geom_col='geometry')
results = [item for item in result]
print(result.shape)
```

**PostGIS : ST\_DWithin**

# Sample Queries

## PostGIS

### Intersects

```
sql = f"""
SELECT *
FROM roads_sydney_500k as roads
WHERE ST_Intersects(roads.geometry,
ST_GeometryFromText('{district_polygon.wkt}', 4326));
"""

start = time.time()
df = gpd.read_postgis(sql, engine, geom_col='geometry')
end = time.time()
time_in_ms = (end - start) * 1000
```

### Within

```
sql = f"""
SELECT *
FROM airbnb
WHERE ST_Within(airbnb.geometry,
ST_GeometryFromText('{district_polygon.wkt}', 4326));
"""

start = time.time()
df = gpd.read_postgis(sql, engine, geom_col='geometry')
end = time.time()
time_in_ms = (end - start) * 1000
```

## MongoDB

```
cursor = db['roads_sydney_500k'].find(
    {
        'geometry': {
            '$geoIntersects': {'$geometry': js['features'][0]['geometry']}
        }
    }
)
print(len([item for item in cursor]), 'documents')
stats = cursor.explain()
time_in_ms = stats['executionStats']['executionTimeMillis']
```

```
cursor = db['airbnb_sydney_1m'].find(
    {
        'geometry': {
            '$geoWithin': {'$geometry': js['features'][0]['geometry']}
        }
    }
)
print(len([item for item in cursor]), 'documents')
stats = cursor.explain()
time_in_ms = stats['executionStats']['executionTimeMillis']
```

# Sample Queries: DE-9IM

PostGIS

MongoDB

Nil